

More About PageRank

Combatting Web Spam

Dealing with Non-Main-Memory Web

Graphs

SimRank

Jeffrey D. Ullman
Stanford University/Infolab



Web Spam

Term Spamming
Link Spamming

What Is Web Spam?

- *Spamming* = any deliberate action intended solely to boost a Web page's position in search-engine results.
- *Web Spam* = Web pages that are the result of spamming.
- SEO industry might disagree!
 - *SEO* = search engine optimization

Web Spam Taxonomy

- *Boosting* techniques.
 - Techniques for making a Web page appear to be a good response to a search query.
- *Hiding* techniques.
 - Techniques to hide the use of boosting from humans and Web crawlers.

Boosting

- *Term spamming.*
 - Manipulating the text of web pages in order to appear relevant to queries.
- *Link spamming.*
 - Creating link structures that boost PageRank.

Term-Spamming Techniques

- *Repetition* of terms, e.g., “Viagra,” in order to subvert TF.IDF-based rankings.
- *Dumping* = adding large numbers of words to your page.
 - *Example*: run the search query you would like your page to match, and add copies of the top 10 pages.
 - *Example*: add a dictionary, so you match every search query.
 - *Key hiding technique*: words are hidden by giving them the same color as the background.

Link Spam

Design of a Spam Farm

TrustRank

Spam Mass

Link Spam

- PageRank prevents spammers from using term spam to fool a search engine.
 - While spammers can still use the techniques, they cannot get a high-enough PageRank to be in the top 10.
- Spammers now attempt to fool PageRank with *link spam* by creating structures on the Web, called *spam farms*, that increase the PageRank of undeserving pages.

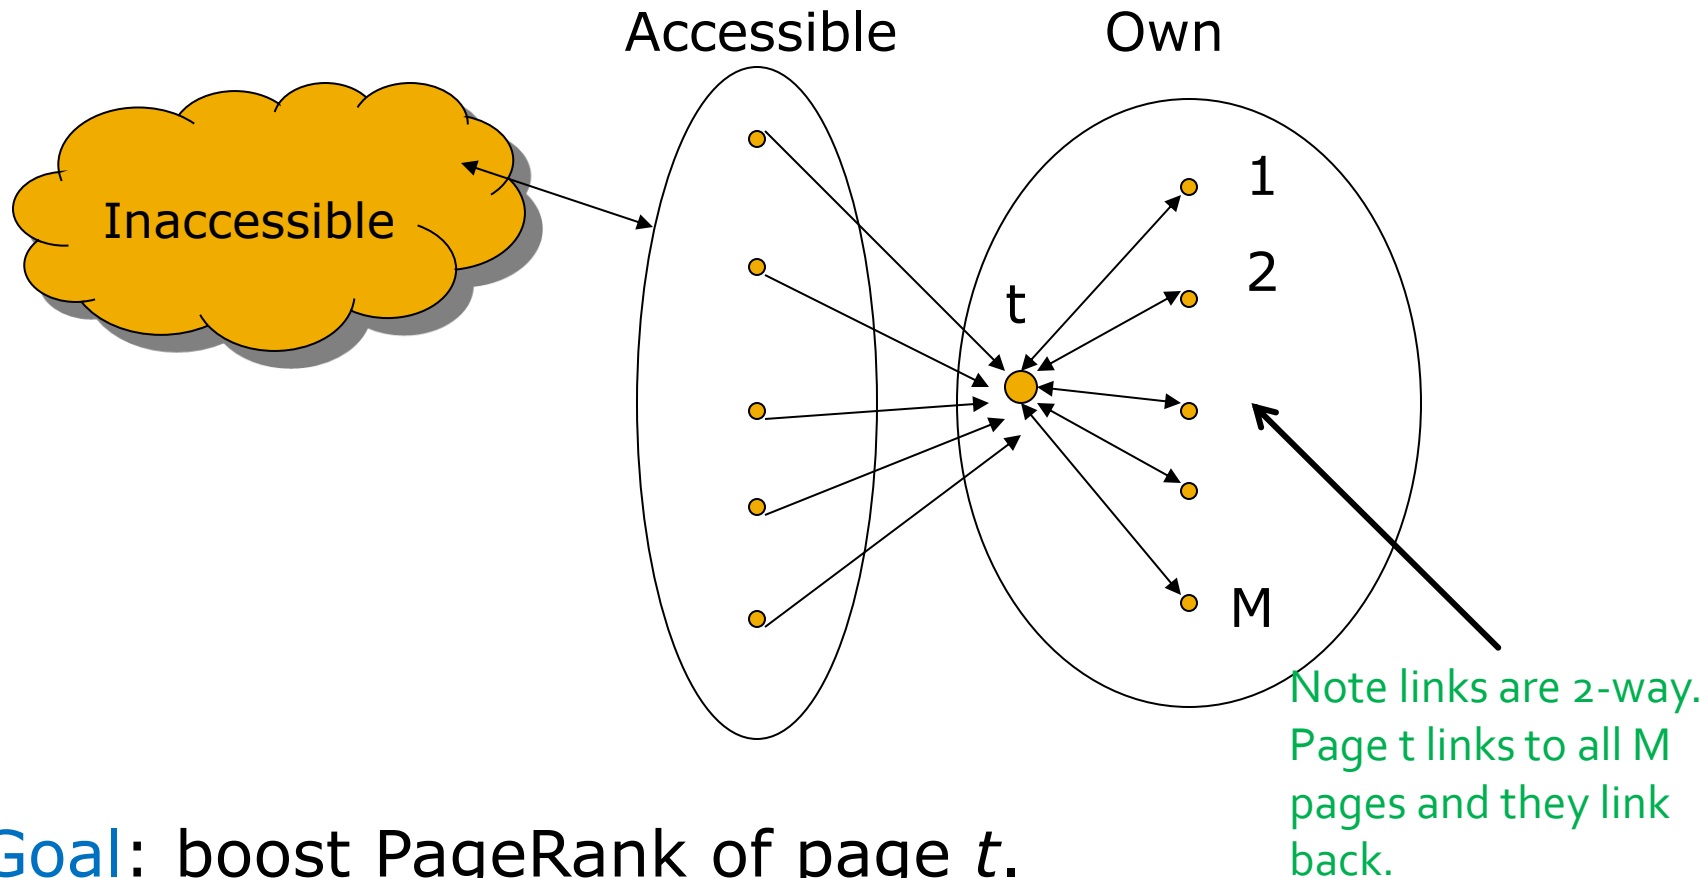
Building a Spam Farm

- Three kinds of Web pages from a spammer's point of view:
 1. *Own pages.*
 - Completely controlled by spammer.
 2. *Accessible pages.*
 - E.g., Web-log comment pages: spammers can post links to their pages.
 - “I totally agree with you. Here's what I wrote about the subject at www.MySpamPage.com.”
 3. *Inaccessible pages.*
 - Everything else.

Spam Farms – (2)


- Spammer's goal:
 - Maximize the PageRank of target page t .
- Technique:
 1. Get as many links as possible from accessible pages to target page t .
 - **Note:** if there are none at all, then search engines will not even be aware of the existence of page t .
 2. Construct a spam farm to get a PageRank-multiplier effect.

Spam Farms – (3)



Goal: boost PageRank of page t .
Here is one of the most common and effective organizations for a spam farm.

Analysis



Suppose rank from accessible pages = x (known).

PageRank of target page = y (unknown).

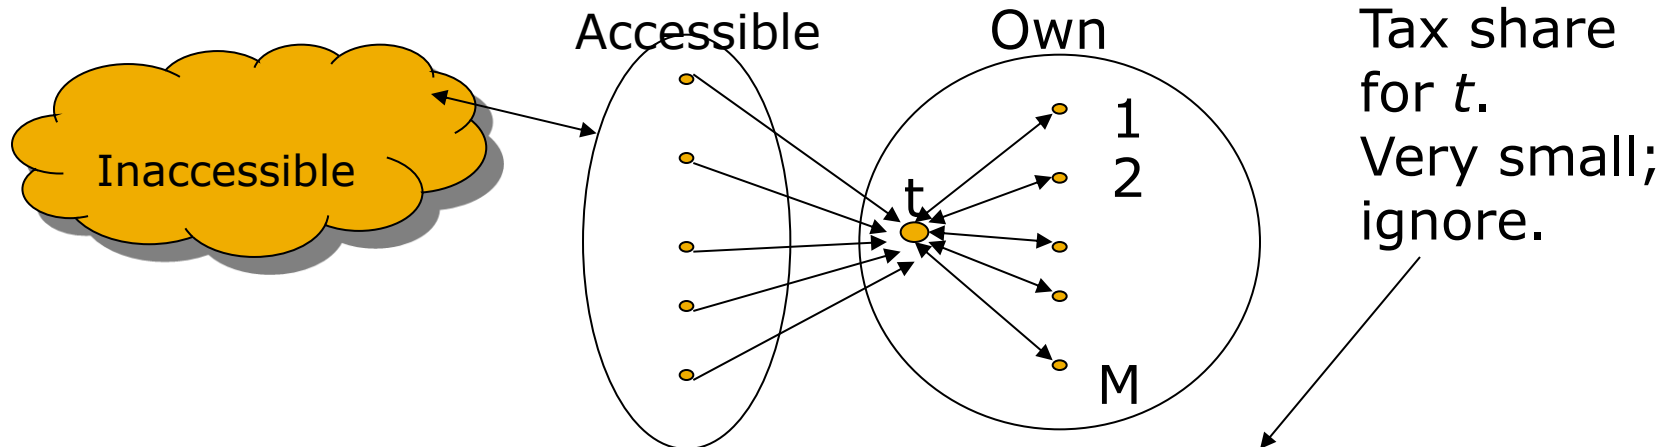
Taxation rate = $1-\beta$.

Rank of each “farm” page = $\beta y/M + (1-\beta)/N$.

From t ; M = number
of farm pages

Share of “tax”;
 N = size of the Web.
Total PageRank = 1.

Analysis – (2)




$$y = x + \beta M [\beta y / M + (1 - \beta) / N] + (1 - \beta) / N$$

$$y = x + \beta^2 y + \beta (1 - \beta) M / N$$

$$y = x / (1 - \beta^2) + c M / N \text{ where } c = \beta / (1 + \beta)$$

PageRank of
each "farm" page

Analysis – (3)



Average page has PageRank $1/N$. c is about $1/2$, so this term gives you $M/2$ times as much PageRank as average.

- $y = x/(1-\beta^2) + cM/N$ where $c = \beta/(1+\beta)$.
- For $\beta = 0.85$, $1/(1-\beta^2) = 3.6$.
 - Multiplier effect for “acquired” page rank.
- By making M large, we can make y almost as large as we want.

Question for Thought:
What if $\beta = 1$ (i.e., no tax)?

War Between Spammers and Search Engines

- If you design your spam farm just as was described, Google will notice it and drop it from the Web.
- More complex designs might be undetected, although SEO innovations are tracked by Google et al.
- Fortunately, there are other techniques for combatting spam that do not rely on direct detection of spam farms.

Detecting Link Spam

- Topic-specific PageRank, with a set of “trusted” pages as the teleport set is called *TrustRank*.
- *Spam Mass* =
 $(\text{PageRank} - \text{TrustRank}) / \text{PageRank}$.
 - High spam mass means most of your PageRank comes from untrusted sources – you may be link-spam.

Picking the Trusted Set

- Two conflicting considerations:
 - Human may have to inspect each trusted page, so this set should be as small as possible.
 - Must ensure every “good page” gets adequate TrustRank, so all good pages should be reachable from the trusted set by short paths.
 - Implies that the trusted set must be geographically diverse, hence large.

Approaches to Picking the Trusted Set

1. Pick the top k pages by PageRank.
 - It is almost impossible to get a spam page to the very top of the PageRank order.
2. Pick the home pages of universities.
 - Domains like .edu are controlled.
 - Notice that both these approaches avoid the requirement for human intervention and (probably) provide adequate distribution.

Efficiency Considerations for PageRank

Multiplication of Huge Vector and
Matrix

Representing Blocks of a Stochastic
Matrix

The Problem

- Google computes the PageRank of a trillion pages (**at least!**).
- The PageRank vector of double-precision reals requires 8 terabytes.
 - And another 8 terabytes for the next estimate of PageRank.

The Problem – (2)

- The matrix of the Web has two special properties:
 1. It is very sparse: the average Web page has about 10 out-links.
 2. Each column has a single value – 1 divided by the number of out-links – that appears wherever that column is not 0.

The Problem – (3)

- **Trick:** for each column, store n = the number of out-links and a list of the rows with nonzero values (which must be $1/n$).
- Thus, the matrix of the Web requires at least

$$(4*1 + 8*10) * 10^{12} = 84 \text{ terabytes.}$$


Integer n

Average 10 links/column,
8 bytes per row number.

The Solution: Blocking

- Divide the current and next PageRank vectors into k *blocks* of equal size.
 - Each block is the components in some consecutive rows.
- Divide the matrix into squares whose sides are the same length as one of the blocks.
- Pick k large enough to fit a block of each vector in main memory at the same time.
 - **Note:** We also need a square of the matrix, but that can be piped through main memory and won't use much memory at any time.

Example: $k = 3$



At one time, we need w_i , v_j , and (a tiny part of) M_{ij} in memory.

Vary v slowest: $w_1 = M_{11} v_1$; $w_2 = M_{21} v_1$; $w_3 = M_{31} v_1$; $w_1 += M_{12} v_2$; $w_2 += M_{22} v_2$; $w_3 += M_{32} v_2$; $w_1 += M_{13} v_3$; $w_2 += M_{23} v_3$; $w_3 += M_{33} v_3$

Representing a Matrix Square

- Each column of a square is represented by:
 1. The number n of nonzero elements in the **entire** column of the matrix (i.e., the total number of out-links for the corresponding Web page).
 2. The list of rows **of that square only** that have nonzero values (which must be $1/n$).
- I.e., for each column, we store n with each of the k squares in one column of the matrix and each out-link with whatever square has the row to which the link goes.

Representing a Square – (2)

- Total space to represent the matrix =

$$\boxed{4*k} + \boxed{8*10} * 10^{12} = 4k + 80 \text{ terabytes.}$$

Integer n for a column is represented in each of k squares.

Possible savings: if a square has all 0's in a column, then n is not Needed for that column.

Average 10 links/column, 8 bytes per row number, spread over k squares.

Note: if 10 is the average number of out-links, then there will be integers in an average of 10 squares for each column, so k can be thought of as the maximum of 10 and the number of blocks.

Needed Modifications

- We are not just multiplying a matrix and a vector.
- We need to multiply the result by a constant to reflect the “taxation.”
- We need to add a constant to each component of the result \mathbf{w} .
- Neither of these changes are hard to do.
 - After computing each block w_i of \mathbf{w} , multiply by β and then add $(1-\beta)/N$ to each component.


Parallelization

- The strategy described can be executed on a single machine.
- But who would want to?
- There is a simple MapReduce algorithm to perform matrix-vector multiplication.
 - But since the matrix is sparse, better to treat it as a relational join.

Parallelization – (2)


- Another approach is to use many jobs, each to multiply a row of matrix squares by the entire \mathbf{v} .
- Use main memory to hold the one block of \mathbf{w} that will be produced.
- Read one block of \mathbf{v} into main memory at a time.
- Read the square of M that needs to multiply the current block of \mathbf{v} , a tiny bit at a time.
- Works as long as k is large enough that two blocks fit in memory.
- M read once; \mathbf{v} read k times, among all the jobs.
 - OK, because M is much larger than \mathbf{v} .

Animation: First Block of v




Main Memory for job i

Animation: Second Block of v



Main Memory for job i

Animation: j-th Block of v



Main Memory for job i


SimRank

Graphs of Entities and Connections
Finding Similar Entities by Random
Walks

Similarity in Networks

- Unlike similarity based on a distance measure, which we discussed with regard to LSH, we may instead wish to look for entities that play similar roles in a complex network.
- **Example:** Nodes represent students and classes; find students with similar interests, classes on similar subjects.

Example: Network

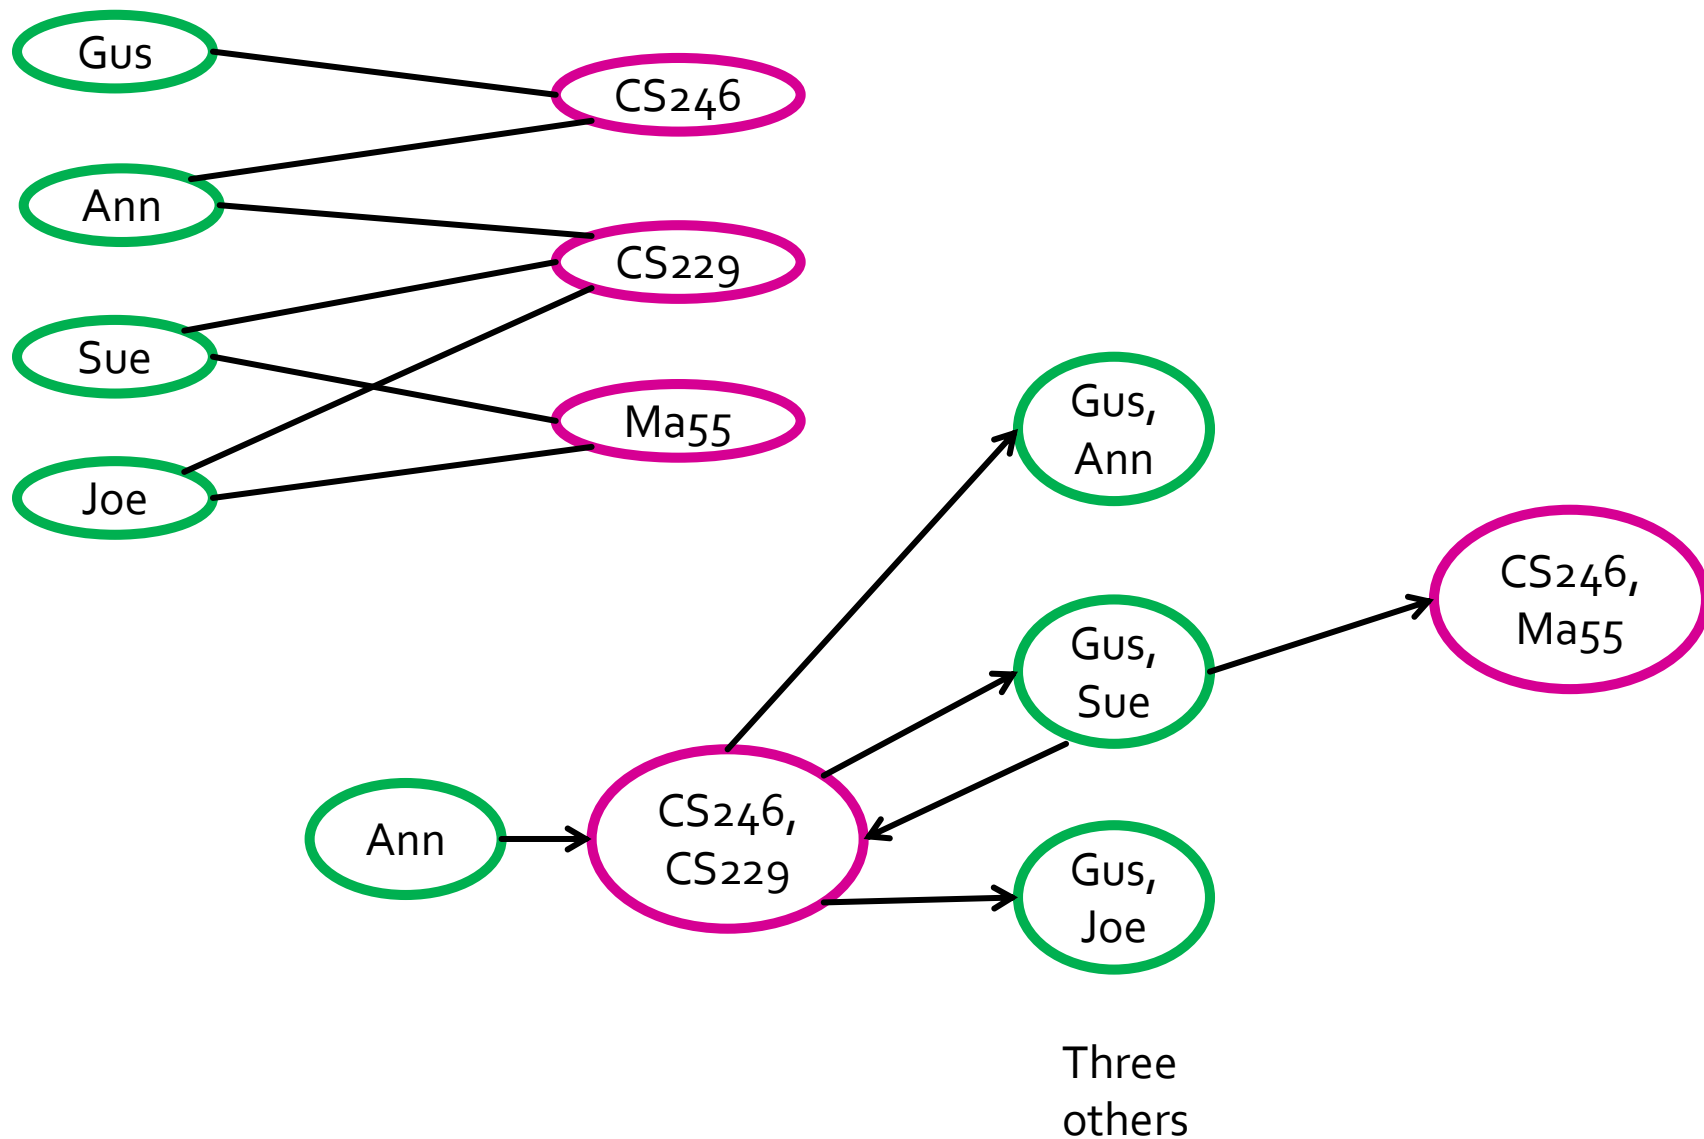


Approach: Pair Graphs

- Intuition:

1. An entity is similar to itself.
2. If two entities A and B are similar, then that is some evidence that entities C and D connected to A and B, respectively, are similar.

Example: Pair Graph




Using Pair Graphs

- You can run Topic-Sensitive PageRank on such a graph, with the nodes representing single entities as the teleport set.
- Resulting PageRank of a node measures how similar the two entities are.
- A high tax rate may be appropriate, or else you conclude things like CS246 is similar to Hist101.
- **Problem:** Using node pairs squares the number of nodes.
 - Can be too large, even for university-sized data.


Alternative: SimRank

- Another approach is to work from the original network.
- Treat undirected edges as arcs or links in both directions.
- Find the entities similar to a single entity, which becomes the sole member of the teleport set.
- **Example:** “Who is similar to Sue?” on next slides.
- Allows us to work on the original graph rather than on pairs.
 - But we need to run many searches (in parallel?).


Example: SimRank




Example: SimRank (20% Tax)



Example: SimRank



Example: SimRank



Example: SimRank

